Partnership proposal

Veterantech × Netra Runtime

Training UPNVJ's next AI engineers on
Indonesian inference

veterantech.id
ksmveterantech@gmail.com
Veterantech
@veterantechupnvj

ABOUT VETERANTECH

Veterantech has shipped 15+ products from UPNVJ's computer science faculty since 2014

Veterantech is a student-led technology organization under the Faculty of Computer Science at UPN "Veteran" Jakarta. We grow members through weekly study clubs, competitions, and products built for real users and clients.

Products developed
15+

Total project valuation
Rp300M+

Operating since
2014

Competition wins
10+

Where our alumni work 30+ ORGANIZATIONS

PROGRAMS AND CORE ACTIVITIES

Five programs take members from weekly classes to client projects

Study Club

A structured weekly learning program across six tracks, each with mentors, a curriculum, and projects. An AI Engineer track would sit here.

Frontend | Backend | UI/UX | Data Science | Android Dev | Product Management

MENTORS FROM

Universitas Indonesia | Telkom University

- AI engineers and PMs in industry

Lombakan

Competition mentoring for national and international tech and business events.

CareerHub

Industry sessions, CV and portfolio workshops, and mock interviews.

Vexpro

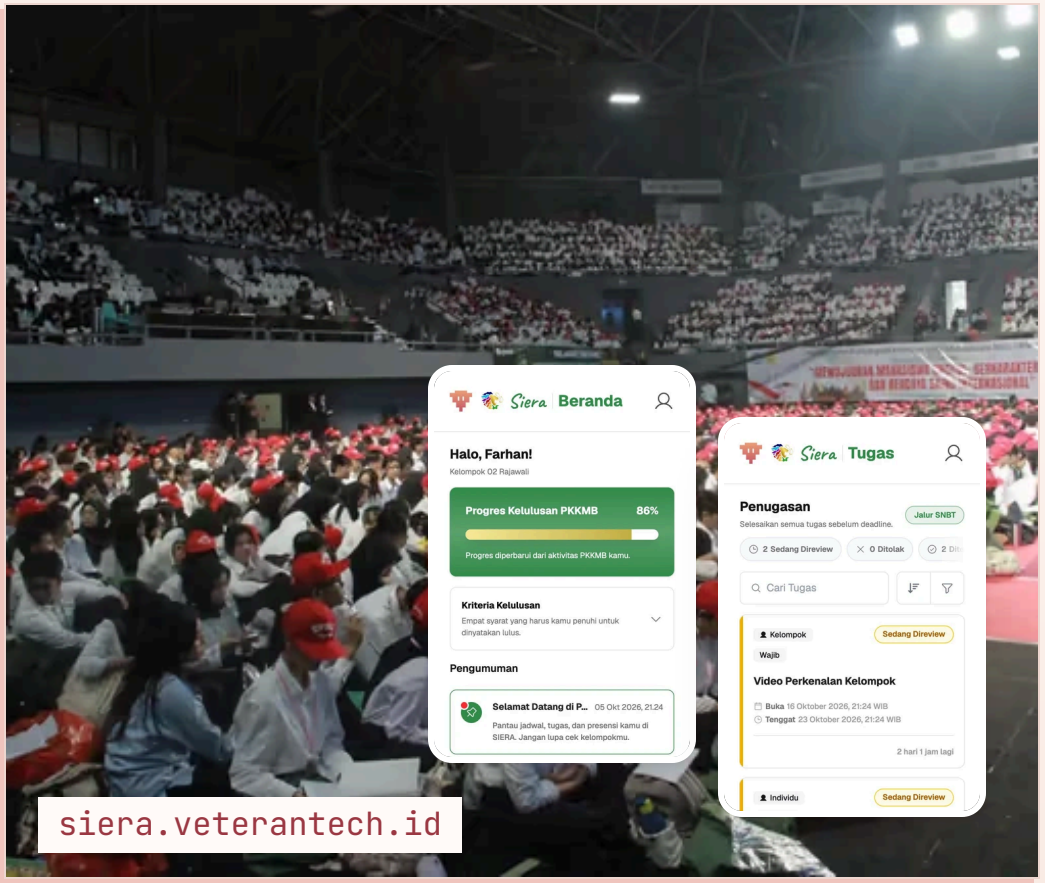
Built for Pesta Warna Nada | FIVE TV FISIP UPN "VETERAN" JAKARTA

VIP

Internal internship across seven divisions of the organization.

PRODUCT · SIERA

SIERA ran UPNVJ's 2026 orientation for 4,900+ users



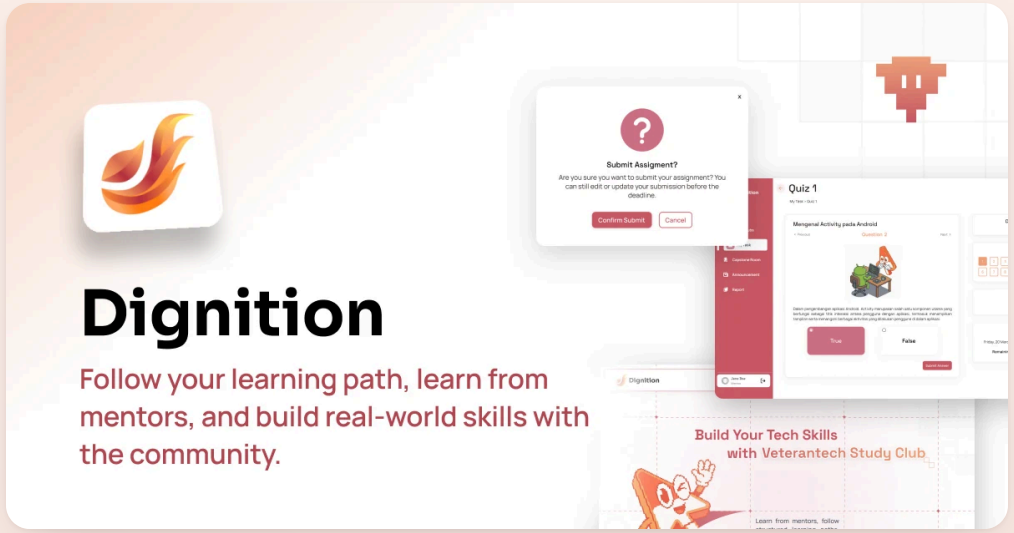
Built in-house by Veterantech for the university. SIERA moved the whole orientation online: automated attendance validation, task management, and real-time, centralized monitoring.

React · Vite · Hono · Redis

Total users	Merch sales	Tracks
4,900+	Rp800M+	5

PRODUCT · DIGNITION

Dignition already runs our Study Club: 740+ learners on our own learning platform



dignition.veterantech.id

Registered learners
740+

Mentors from UI, Tel-U, and industry
22

Study clubs
19

Learning resources
217

Our own LMS for mentoring, assignments, and progress. Its MCP connector already lets mentors and mentees use it from Claude, ChatGPT, and Cursor.

WHERE NETRA COMES IN

Our next track needs an inference partner

We have used Netra at NalarX. The pilot puts a Veterantech class on it and measures how it performs.

01	What we saw using Netra at NalarX	07
02	The AI Engineer pilot	08
03	Abuse control and measurement	09
04	What we commit to	10



VETERANTECH MEMBERS · UPNVJ

NETRA AT NALARX

From using Netra at NalarX: easy to build on, and throughput is what we watch

WHY IT FITS A CLASS

Netra Runtime

- **OpenAI-compatible API**
  Existing SDK code only changes its base URL
- **Keys you can control**
  Expiry, RPM/TPM limits, and budgets in the dashboard
- **Data stays in Indonesia**
  Zero data retention by default, which suits a state university

WHAT WE WATCH: OUTPUT SPEED

tokens/sec, typical	tokens/sec, frequent dips
~400	<200

Seen while building NalarX, not a benchmark. One gateway in front of 24–30 builders turns this into numbers Netra can act on.

THE PILOT · PROPOSED

An 8-week AI Engineer track: 8 teams ship live AI apps on Netra

WEEK	MODULE	MILESTONE
1	Setup and persona: Netra through the gateway, streaming, errors	Gateway connected
2	Cost and routing: structured JSON, semantic cache, model router	Assignment 1 (solo)
3	Tools and MCP: function calling, agent loop, local MCP server	Tools running
4	Vectors and RBAC: embeddings, vector DB, role filters	Role-based search
5	Secure RAG: retrieval quality, citations, RAG evaluation	Assignment 2 (team)
6	Guardrails: prompt injection, PII redaction, red-teaming	Attacks measured
7	Observability and deployment: tracing, metrics, Docker, HTTPS	App live on a server
8	**Demo day with Netra**	**8 apps presented**

Sizing basis: our academic team's plan: 24–30 mentees selected from Backend and Data Science alumni, 8 teams of 3–4.

ABUSE CONTROL AND MEASUREMENT

Mentees never hold a Netra key: every request passes a gateway that caps and measures it

01 · MENTEE

Mentee's code

Standard OpenAI SDK, our base URL, a personal key.

02 · VETERANTECH

Our LLM gateway

- Hard US$ budget per key
- Rate limits per key
- Only the models we allow
- Tokens, cost, and timing per request

03 · NETRA

Netra Runtime

Netra API

One real key, held by us only.

One call revokes one key; nobody else is affected.

Built and tested: budget block, key revoke, cost logs.

WHAT WE COMMIT TO

Netra gets product feedback from 8 live apps, plus public showcases and a path into client work

WE COMMIT TO	WHAT NETRA CAN DO WITH IT	WHEN
Two usage and feedback reports	**Who builds, where they got stuck, structured output and tool-calling quality, and speed at peak load**	**Weeks 4 and 8**
8 open-source apps	Public repositories with documentation, a short tutorial, and Netra credited as the inference provider	Week 8
A showcase that tags Netra	Final projects featured on Veterantech's LinkedIn and social channels, with Netra tagged	After demo day
Netra first for client AI work	When a Vexpro client project needs LLM inference, we propose Netra first	After the pilot

HOW WE MEASURE

Netra sees what its own dashboards can't: who builds, where they get stuck, and speed from the user's side

	WHAT WE REPORT	SOURCE
ADOPTION	Weekly active builders and tokens	Gateway logs
	Apps live by demo day	Public repos
FEEDBACK	Where new developers get stuck: docs, errors, SDK gaps	Mentor log
SPEED	Time to first token and tokens/sec, p50 and p95	Gateway logs
	Error rate on Netra's side	Gateway logs

WEEK 1	Keys issued, gateway connected, kickoff
WEEK 4	Interim report and review with Netra: adjust quotas, models, or scope
WEEK 8	Demo day and final report
NEXT	Decide on a second cohort or a hackathon

THE ASK

Three yeses today let us start the first cohort

01 Inference credits for the pilot
About 100M tokens over 2 months, priced on the models we agree. An education credit program works too.

02 Endpoints the class depends on
Tool calling, an embedding model, bursts up to 20 RPM, and 7 days' notice before model or price changes.

03 Netra in the class
A guest session by a Netra engineer and a seat on the demo-day panel, online or on site.

Credit basis: 30 mentees × 1.5M tokens (weeks 1–4) + 8 teams × 5M tokens (weeks 5–8) + 15M buffer for testing and demo day.

× Netra Runtime

EMAIL ksmveterantech@gmail.com
LINKEDIN Veterantech
INSTAGRAM @veterantechupnvj